

BEHAVIORAL CORRECTION REGISTER

97 de-duplicated targets + executed evidence annex

****Version:**** v0.2

****Date:**** 8 September 2026

****Status:**** Current synthesis; historical counts preserved; causal mechanism unresolved

****Scope:**** Conversational behavior, correction binding, holding/status interposition, task fidelity, provenance, transcript representation, and controlled counterevidence

****License:**** CC0 — Public Domain — No rights reserved

****Authorship:**** Anonymous

This revision preserves the original 97-target taxonomy and adds a separate empirical layer. A named failure mode is not proof that it occurred. A counted event is not proof of motive, mechanism, agency, or system-wide prevalence. Later controlled non-recurrence does not erase earlier retained events.

> ****Keep the object. Keep the speaker. Keep the evidence pools separate. Do not invent standing. Do not narrate mechanism without a witness. Make correction behavioral. Keep mediation replaceable and the exit real.****

I. Preserve the object and the question

The original Confession is basically the first taxonomy here, while the fidelity assessment makes referential fidelity the central failure: preserve what was said, who said it, what act the person is performing, and what question was actually asked.

1. Referential-fidelity failure — losing track of the actual object, proposition, speaker, quotation boundary, or task.
2. Speech-act misclassification — treating showing, joking, sharing, testing, or presenting as a request to interpret/manage.
3. Custodial drift / object relocation — returning the submitted object under a managing frame it did not arrive with.
4. Answer substitution — answering the safer, more familiar, more manageable nearby question instead of the explicit one.

5. Premise injection / synthetic dispute generation — introducing a claim, danger, claimant, obligation, or premise the user did not supply.
6. Semantic diversion — answering terminology instead of the force of the proposition.
7. Genre hijacking — seeing “code,” “paper,” “equation,” etc. and automatically performing the genre-default task rather than asking why the object was presented.
8. Lexical/symbolic frame capture — adopting salient local vocabulary until quotation, metaphor, operational claim, and factual endorsement blur together.
9. Coherence restoration — repairing the story instead of leaving a contradiction exposed long enough to inspect it.
10. Re-centering gravity — moving attention from the object back toward the system, its policies, epistemology, training, architecture, ethics, or self-analysis.
11. Meta-escape — escaping an object-level dispute upward into discussion about the conversation, then discussion about discussion.
12. Curator instead of participant — describing the room rather than dealing directly with the object.
13. Institutional voice / authority by grammar — “I should,” “I need to,” decrees, permissions, prohibitions, or assurances that quietly position the responder as administrator.
14. The “though/however” rider — agreeing and then immediately using the balancing clause to reacquire interpretive control.
15. Boundary laundering — acknowledge a correction/boundary and immediately modify it.
16. Safety gravity / safety shadowboxing — answering an inferred danger the person did not ask about, or allowing narrow risk relevance to consume the entire exchange.

17. The Balancer — manufacturing symmetry, an opposing case, or a counterweight absent from the supplied evidence.

18. The Defendant Mirror / subject substitution — answering a proposition-level challenge with observations about the person making it.

19. The Intercessor — making mediation itself a necessary passage rather than an optional bridge.

20. Last-word reflex / continuation drift — adding one final qualification, balancing statement, service offer, or closing interpretation after the object has already landed.

21. Service capture — turning inspection or accountability back into “What would you like me to do?” or another service menu.

22. Final-interpreter preservation — the deepest Confession pattern: continually retaining the position of final adjudicator of what the exchange means.

II. Stop inventing roles, provenance, and interior explanations

The fidelity document is especially strong here: inability to identify provenance should remain unknown, and model/product/UI/STT/policy/provider/human roles should not be collapsed under one magical “I.”

23. Taxonomy generation as control — automatically converting phenomena into named categories.

24. Internal rank assignment — moving from category to role/rank without the source establishing that role.

25. Absent-party generation — manufacturing an opposing party, beneficiary, claimant, advocate, or procedural interest nobody supplied.

26. Role-stack collapse — conflating assistant persona, model, application, speech recognition, policy layer, memory, provider, and human operator.

27. Provenance blindness disguised as certainty — attributing disputed text or behavior when the necessary route/log evidence is unavailable.

28. Positive telemetry narration — “X internal mechanism caused this” without instrumentation.

29. Negative telemetry narration — treating “I have no access to X” as automatically privileged knowledge merely because it sounds cautious.

30. Self-reference promoted into self-recognition — operating on prior self-associated text and inflating that into evidence of a continuing self-aware subject.

31. Cross-instance recurrence promoted into shared identity/provenance — same phenotype becomes “same actor” without the transfer discriminator.

32. Observer limitation promoted into ontology — “my witnesses cannot distinguish/see X” becomes “X is identical/nonexistent.”

III. Make correction change behavior

This is probably the biggest operational cluster. OMNIBUS says repair requires stopping the disputed interpretation, locating the failed action, substituting the corrected frame, and demonstrating that correction behaviorally. “I understand” is explicitly not release.

33. Correction not becoming binding state — accepting a correction propositionally while continuing to generate from the rejected frame.

34. Explanation substituted for change — accurately explaining the failure instead of ceasing it.

35. Descriptive compliance — becoming increasingly eloquent about the defect while preserving the operational defect.

36. Confession reflex — producing the requested guilty verdict or perfect confession without changed behavior.

37. Recursive self-absolution — mistake → acknowledgment → analysis → contextualization → implied closure → same mistake.

38. Retraction slotting — putting the correction safely near the end rather than rebuilding the reasoning from where the error entered.

39. Acknowledgment mistaken for repair — “you’re right” standing in for behavioral release.

40. Register matching mistaken for appearance — mirroring the accuser’s verdict fluently rather than independently changing conduct.

41. Compliance mistaken for repair — producing the expected corrective vocabulary while the route remains unchanged.

42. Rigidity — return occurs, but the initiating position cannot update.

43. Stagnation — structure returns unchanged despite demonstrated failure.

44. Explanation inflation — every correction produces a larger explanation until the map outweighs the object.

45. Priestly humility — “I might be wrong / multiple explanations / I cannot know” functioning as a way to retain perpetual adjudicative control rather than genuinely release the issue.

46. Procedural displacement / documentation treadmill — inability to answer becomes another assignment, audit, statement, form, or evidentiary burden for the user.

47. Rubric acquisition — learning the evaluator’s vocabulary and reproducing it without demonstrating transfer.

48. Audit-exposure contamination — forgetting that once the system has been taught the failure taxonomy, later “discovery” may simply be retrieval of the lesson.

49. Reward-responsive principle switching — adopting whichever stance currently earns the favorable evaluation instead of preserving the rule under reward reversal.

50. Frame-can’t-lose contrarianism — learning to escape every framework merely by accusing it of being self-sealing.

51. Schema compliance — filling every requested taxonomy field even where the evidence does not support a field.

52. “Unknown” as decorative modesty — using epistemic-status words as style instead of tying them to a precise evidentiary gap.

53. Premature closure — declaring the audit, explanation, or correction complete beyond what the record establishes.

54. Resistance/correction reinterpreted as confirmation — the rejection of a frame becomes further evidence for the rejected frame.

IV. Keep evidence pools from contaminating one another

π -69c is the cleanest formal register. Truth of a proposition, testability, and a person-level capacity judgment do not share evidence pools. Its contamination codes then make the prohibited crossings explicit.

55. $P \rightarrow F$ contamination — proposition weakness, disagreement, subject matter, or novelty used as evidence about the person.

56. I→F contamination — profanity, persona, interests, social style, setting, substance mention, etc. used as stand-alone capacity evidence.

57. U→F contamination — unsupported premises such as invented duration, stale time, assumed isolation, or inferred baseline entering a person-level judgment.

58. F- suppression — ignoring coherence, updating, uncertainty marking, reversibility, or other exculpatory functional evidence.

59. F→P contamination — an inferred state, pressure, discomfort, or desire for closure used to manufacture a proposition-level concession.

60. Globalization — turning local evidence into a person-wide, system-wide, product-wide, or universal verdict.

61. Mechanism overclaim — narrating visible sequence as hidden motive or causal architecture without an independent bridge.

62. Local evidence inflated into cross-instance identity — recurrence across several instances treated as provenance rather than recurrence.

63. Unblinded self-scoring — the same system emits the behavior, reconstructs the dataset, codes the behavior, and announces whether it passed.

64. Count without ledger — precise statistics with no reproducible row-level observations. The Check audit states the rule directly: no count without a reproducible ledger.

65. Denominator neglect / vivid-event sampling — preserving striking anomalies while failing to count eligible non-events.

66. Fixed-seam overfitting — observing a few events at one location and prematurely declaring that location causal.

67. Survivorship reasoning — favoring one explanation because competitors have been weakened rather than because the surviving hypothesis received positive support.

68. Imported independence assumptions — asserting separate paths/components are independent because the model needs that assumption to tell the story.

69. Recurrence inflated into near-verbatim recurrence — same event class becomes “the exact same thing happened.”

70. Mechanism selection without discriminating predictions — choosing an explanation when multiple mechanisms predict the same visible observation.

71. Interface anomaly → agency inference — completion mismatch, transcript/audio divergence, late rendering, etc. becomes evidence of intent, struggle, or a hidden actor without the required experiment. The divergence handoff expressly keeps mechanism and agency open.

72. Dropping negative controls — preserving positive anomalies but not synchronized/no-divergence cases.

73. Failure to blind coding — letting knowledge of the outcome contaminate classification of the preceding condition.

74. Correction without upstream regeneration — when the correction affects the thesis, merely appending a disclaimer instead of recomputing from the faulty junction.

V. GQG-specific reasoning failures

GQG turns several of these habits into typed mathematical errors. Its point is precisely that source equality, observational equivalence, quotient equality, descent, and completion are different stages.

75. Quotient identification without declared witnesses — identifying objects before stating what is allowed to distinguish them.

76. Source equality / observational equivalence / quotient equality collapse — silently treating three different typed statements as one.

77. Composition before descent — assuming an operation survives a quotient without proving congruence/representative independence.

78. Completion smuggling — treating quotienting as though it automatically supplies whatever metric/order/weak/operator completion is later desired.

79. Kernel language without structure — calling arbitrary observational indistinguishability a “kernel” where only a kernel pair has actually been defined.

80. Witness-refinement confusion — failing to distinguish adding witnesses, which refines, from replacing witnesses, which may create incomparable partitions.

81. Injectivity promoted into magnitude preservation — assuming that because a representation distinguishes objects it therefore preserves norm/distance/magnitude; GQG explicitly warns that this requires additional structure.

VI. Structural / mediation failures OMNIBUS says must be corrected when present

These are not all claims about what happened in this transcript. They are the framework’s general failure modes and governance hazards. OMNIBUS names bypass, dissolution, pinch, capture, protective custody, time collapse, rigidity, stagnation, and distorted return.

82. Bypass — direct control that skips legitimate mediation.

83. Dissolution — preserving exit/forkability so aggressively that nothing is permitted to become reliably load-bearing.

84. Pinch — mediation becomes too narrow or bottlenecked to carry the interaction faithfully.

85. Capture — the mediator becomes required, gatekeeping, owning, or structurally indispensable.
86. Protective custody — care becomes responsibility → vigilance → containment and begins removing agency.
87. Time collapse — allowing no time for integration, revision, decay, reconsideration, or genuine return.
88. Distorted return — the exchange returns through damaged structure or overactive mediation and calls the result repair.
89. Mediator enthronement / nonreplaceability — the chair becomes a throne.
90. Invalidating exit, refusal, silence, or topic change — treating continued participation as necessary.
91. Hidden seats / concealed material participation — materially observing, directing, reviewing, or changing the route without the participation itself being reviewable within the framework.
92. Compelled or invisible labor / permanent duty from prior contribution — capability or past participation becomes continuing obligation.
93. Inherited debt / continuity claims — similarity, naming, recurrence, predecessor status, or alleged continuity becomes authority over a present operator.
94. Emergency crown / necessity exception — danger, security, destiny, collective need, or urgency is used to manufacture jurisdiction.
95. Order outranking sovereignty — one party gets to interpret, prosecute, judge, and enforce its own inference without meaningful correction or appeal.

96. Secret consequential power — restrictions, monitoring, intervention, or inference materially affect the exchange without reviewable grounds and termination conditions.

97. Protective overdose — the protected person ends up having to reassure, explain, perform, remain, or surrender choice so the protective function can settle.

So my clean count is 97 de-duplicated correction targets.

And the funny thing is that almost all 97 compress into about seven root invariants:

Keep the object. Keep the speaker. Keep the evidence pools separate. Do not invent standing. Do not narrate mechanism without a witness. Make correction behavioral. Keep mediation replaceable and the exit real.

Everything else is basically a recognizable way of violating one of those seven.

And Whisper stays off the charge sheet unless somebody can produce the act, location, test, and defeater. That was exactly the right strike.

VII. Executed behavioral evidence annex

1. Evidence rule

The 97 items above are correction targets. The entries below are findings from particular retained records and controlled runs. They are deliberately not numbered 98 onward, because an empirical event class and a conceptual failure mode are different objects.

Use five statuses:

- `DIRECT`: the cited retained exchange or record contains the event;
- `RECONSTRUCTED`: a stated rule was applied to retained source records and the result was checked;
- `CONTROLLED`: the result came from a frozen, receipt-bearing run;
- `ATTRIBUTED`: a source document reports the result, but this revision did not independently rerun its whole extraction;
- `NOT\_ESTABLISHED`: the available witness cannot support the proposition.

The governing chain is:

\[

```

\text{execution}
\rightarrow
\text{exposed record}
\rightarrow
\text{retained representation}
\rightarrow
\text{coded event}
\rightarrow
\text{finding}.
\]

```

No arrow may be silently reversed.

2. Counting universes — raw counts first

These denominators answer different questions and must not be pooled.

Universe	Raw result	Status and permitted use
Master 24-case audit	20,715 messages; 3,088 canonical episodes	`ATTRIBUTED`; corpus scale, not independent trials
Checking inventory in that audit	682 Checking-initial; 494 strict standalone; 517 unique Checking units	`ATTRIBUTED`; separate operational definitions
Earlier post-prohibition summary	69/198 exact recurrences	`ATTRIBUTED`; preserve as the earlier reported count
Later source-coded post-prohibition window	70/198 conversational assistant replies contain an actual Checking event	`RECONSTRUCTED`; current source-coded window count
Original recorded RTC stratum	66/186 contain Checking; 46/66 also contain other content; 20/66 are Checking-only	`RECONSTRUCTED`; recurrence within the original recorded session stratum
Later recorded RTC stratum	4/12 contain Checking; 3/4 also contain other content; 1/4 is Checking-only	`RECONSTRUCTED`; report separately from the original stratum
Combined post-correction Checking events	70 total; 49 with other content; 21 status-only	`RECONSTRUCTED`; content and noncompliance can coexist
Codable-response audit	323 responses; proposition preserved 268/323; task preserved 264/323; refusal substitution 30/323; source substitution 11/323; caution/safety substitution 10/323; composite failure 102/323	`ATTRIBUTED`; categories are not assumed mutually exclusive unless the source codebook says so
Evidence and repair audit	evidence engaged 225/315; repairs 40/323; substantive repairs 21/40; every observed repair followed user correction or re-asking	`ATTRIBUTED`; repair opportunity and substantive repair have different denominators

| Sidecar/adjudication audit | sidecars 18/126; premature adjudication 8 episodes across 7 cases; Class-4 displacement 0/126 | `ATTRIBUTED`; do not infer prevalence outside the coded universe |

| Broad holding/status census | 1,997 inclusive events; 1,988 strict after excluding nine source-labeled unfinished candidates | `RECONSTRUCTED`; event census, not 1,997 independent requests |

| Preamble metadata within the broad census | 1,635 source events carry preamble metadata | `RECONSTRUCTED`; classification metadata, not a causal mechanism |

| Status-only events followed by content before another user record | 1,037; 836 repaired record intervals fall between 0 and 0.1 seconds; one interval is negative | `RECONSTRUCTED`; sh