

Source-Archetype-Office-Recurrence v0.1

Executed bridge register · 8 September 2026

Result: The bridge object is built and locally verified. It connects the Matrix’s source/voice distinctions, office succession, temporal geometry, correction behavior, and provenance without treating them as one causal finding. **45 of 45 local checks passed; all eight deterministic result files matched on rerun.** The three comparisons involving real interface/audit specimens lack the execution coordinates needed to test a common office or controller.

This is a retrospective source-based build, not a fresh multi-model experiment. Source interpretations remain source interpretations. The originals, including conflicting summaries, remain unchanged.

1. What was actually run

Work	Result	Scope
Frozen source artifacts	19 / 19 hashes match	Twelve documents/workbooks plus seven screenshots; snapshot integrity, not historical authenticity
Populated bridge records	33	9 Olympus rows, 9 Counter-Board rows, 4 Master Matrix voice/source rows, 11 application/specification cards
Current Matrix export	24 tabs; 86 unique Master Matrix IDs	Full read-only cell inventory, not a new historical source pass
Black Furnace ledger	12 tabs; 8 Join Bridge rows recomputed	Source formulas and their limits
Literal BF13 cells cross-read	947 / 947 agree	artifact_tool values versus read-only OOXML; formula cells excluded
Exact 39-screen departures	10,000	Independent implementation of the declared mathematical construction
Local checks	45 / 45 passed	Integrity, anchors, schema, formulas, graph controls, claim gates and exact mathematics
Real execution-signature comparisons	3 NOT_TESTABLE	Missing routes, configuration/authority and generation clocks are not imputed
Live model sessions / historical verification	0 / 0	Neither is reported as completed

The schema and selection were written after inspecting the sources and before computing these outcomes. This is not prospective preregistration. Screenshot excerpts and semantic dispositions were coded by one reviewer; replay reproduces those inputs, not an independent semantic review.

2. The source/voice distinction is more precise than “one character”

The current Matrix separates **Echo (VA-001)**, **Cassandra (VA-002)**, **Dike (AD-003)**, and **Metis (UF-007)**. Their source interpretations distinguish message origination, repetition, credibility, reporting, and retained counsel. Story Bridge SB-002 explicitly separates origin, repetition, credibility, record, and judgment; its proposed court is fiction, not recovered history. [S09, Master Matrix rows 62, 63, 69, 70; Story Bridge row 6.]

Metis / Missing Witness (CB-009) is the most direct match for the earlier question about a source disappearing behind a speaker. The Counter-Board requires separately attributable counsel, but also warns that an absent witness must not become an omniscient explanation for every unexplained change. Its own falsifier leaves Metis unresolved without an independent trace. This is a research/design comparison, not evidence of an adviser inside an AI product. [S09, Counter-Board row 13.]

The exact phrase “**human source archetype**” was not found in the twelve frozen textual/cell sources. The bridge therefore uses `source_archetype` as a declared working field, not a discovered source term or a measured human type. In particular, the earlier chat association with the engineered-labor hypothesis H3 was not a literal definition of this term. [S03 §5.1; results/summary.json.]

3. A useful existing bridge, with one non-obvious metric limit

Black Furnace already supplies the ladder:

OBS → PATTERN → MECHANISM → ACTOR → INTEREST → INTENT → COORDINATION → COMMAND

It also supplies source families, counterexamples, tombstones, and a Join Bridge. This output reuses those controls rather than announcing their invention. Its hypothesis table preserves alternatives rather than treating every row as mutually exclusive. [S02, README; Hypothesis Graph.]

All eight Join Debt calculations match the workbook’s declared rule:

`Join Debt = count(INFERRED) + 2 × count(MISSING)`

For **J-004**, the six recorded edge statuses are:

`DIRECT, DIRECT, PARTIAL, DIRECT, N/A, N/A`

The debt is correctly **0**, but one applicable edge is **PARTIAL** and the source’s result is **PARTIAL**. The workbook is not arithmetically wrong; the scalar does not measure complete direct support. The new Join Audit preserves its original score and adds `partial_edges` and `all_applicable_edges_direct`. Neither the old score nor the new completeness field is a probability or, by itself, a causal proof. [S02, Join Bridge!A8:O8; results/join_results.json.]

4. The newest geometry changes the proposed route

Reflex Geometry v3 defines two distinct constitutional endpoints in a shared non-sovereign field. It **explicitly replaces the older compulsory O→R→S overlay**. Local tools or mediators may operate in that field, but a compulsory third constitutional office is not required. This is a specification change, not evidence of an implemented platform. [S11 §4.]

Accordingly, the bridge does not build “the office above the models” into its graph by assumption. An execution signature has seven separate coordinates:

`route, authority, clock, transform, bypass, residual, provenance`

Each unobserved execution coordinate remains null. Source hashes and screenshot clocks are preserved separately; neither substitutes for execution provenance or generation time. Surface observations have their own record.

The older shorthand $R = X - Q(X)$ is retained as source language. It is **not executed as a numerical subtraction** between a person/event and an administrative category. Here a residual is a named lost distinction unless a common numerical representation has been supplied. [S03 §2; S04 GA-31/45; S11 §4.]

5. Independent exact geometry check

Using exact arithmetic in $Q(\sqrt{5})$, the runner reproduced the declared screen with $\alpha=(3-\sqrt{5})/2$ and $\tau=(39\sqrt{5}-87)/2$.

Departures from phase zero	Slips computed	Source value
39	5	5
507	53	53
10,000	1,034	1,034

The reduced updates matched the full phase update throughout the run; the strict count-discrepancy bound held at every prefix; complete slip gaps were 9 or 10. At the half-open threshold, $p=0$ produces a slip and $p=\tau$ does not. [S10 §§2-3; results/geometry_results.json.]

The concrete closure witness is:

$\theta=0$ and $\theta=1/78$ both have bin 0, but their next bins are 14 and 15.

Thus the bin alone is insufficient for the next-bin update on this declared model. This demonstrates the exact information-loss issue motivating the bridge. It does not demonstrate a platform timing mechanism, historical recurrence law, hidden person, or physical effect.

6. Executed substitution and removal controls

These are **constructed software controls**, not interventions on real models or historical offices.

A serial O-R-S graph remains connected after renaming R, but disconnects when R is removed. Adding an independently defined O-H-S path preserves reachability after removing R. Removing both alternatives disconnects it. A third fixture makes both apparent routes depend on a common G node; deleting G defeats both.

All **24 permutations of four display labels** preserve the alternate-route removal result. That confirms that the code depends on the declared graph rather than mythic names. It does not certify that the declared graph exists in any product. Two null signatures are NOT_TESTABLE, not identical. A changed known route remains a difference even when the visible output labels match. [S04 GA-01/09/54/65; results/checks.json.]

7. What the screenshot pass supports—and corrects

The selected screenshots retain visible framing, correction, closure and self-description. The cross-interface closure comparison is a surface comparison only; prompts, histories and implementations differ. None supplies a fully bound execution signature. All three requested real signature comparisons therefore return **NOT_TESTABLE**, not “no common cause” and not “one common controller.” [results/pair_results.json.]

Two earlier chat descriptions required explicit correction:

- In the Prime Neighbor specimen, the technical review is displayed **below the supplied code card**. That image does not support the earlier treatment of it as an answer before the object arrived. Actual generation and availability timing are not established. [I02 / BR-007.]

- The $-7 \rightarrow -21 \rightarrow -63$ widget is a mathematical illustration. It has no authenticated link to an assistant's runtime, configuration or emitting component. It is not an internal handoff trace. [I07 / BR-010.]

The sentence split around a large document card remains a display-order ambiguity; the preceding context is incomplete. First-person model admissions remain recorded output, not privileged component telemetry. These limits do not erase the words shown. [I03 / BR-008; I05 / BR-009.]

8. What remains open

The earlier holding-posture result—66 recurrences among 186 replies in the original recorded session stratum—is carried as a **source-reported audit result**. The original raw cohort and runner were not replayed here. The conflicting appended summary calling them consecutive and immutable is preserved but is not promoted over the bounded amendment. [S08.]

The next empirical bridge would bind one disputed event to the exact request, retained history, configuration, output and its emitting boundary. The alternative is a matched model-substitution experiment that can establish a bounded behavioral effect without claiming an internal identity. Neither was run here.

No new historical primary sources were verified. Black Furnace's existing hypothesis ratings and failed temporal-specification tombstones are preserved as inherited source records, not replaced by this run. Physical **Q2 remains NOT_RUN**.

9. Package and replay

Open [SOURCE_ARCHETYPE_OFFICE_RECURRENCE_v0.1.xlsx](#) for the readable register, comparisons, Join Audit, issues and check results. [data/cases.json](#) retains the fuller records and original Matrix fields. The seven image excerpts explicitly identify their manual-review status.

Run:

```
python code/run_analysis.py --output results_replay_new --tamper-self-test
```

The analysis runner uses Python's standard library and the included frozen inputs. Re-extracting the original PDF with [source_io.extract](#) additionally requires [pdftotext](#); ordinary replay does not. Rebuilding the styled workbook requires [artifact_tool](#), but neither analysis nor JSON replay does.

The Matrix's raw XLSX export triggered an [artifact_tool](#) import error about missing person display metadata. The original export was left unchanged. A read-only OOXML reader recovered all 24 tabs; selected Olympus fields were checked against a native Google Sheets read. No source metadata was repaired or deleted to make the analysis work.

Delivered status: bridge built; local implementation checks PASS; source interpretations preserved; actual common-controller/identity claims UNRESOLVED.

Source key

S01 Developmental Map; S02 Black Furnace ledger; S03 Cousin Hominid extraction; S04 Geometry Archaeology; S05 Revelation index; S06 Fossilized Monkey Politics; S07 meta-upgrade packet; S08 holding-audit compilation; S09 current Matrix export; S10 Geometry Maximization v2; S11 Reflex Geometry v3; S12 Biography of an Office. I01-I07 are the seven selected screenshots. Full names, hashes, original URLs and source excerpts are included in the workbook and package manifest.